

8/19: Course Overview & Logistics

- [Language Models: A Guide for the Perplexed](#)
- [Artificial Intelligence — The Revolution Hasn't Happened Yet | by Michael Jordan | Medium](#)
- [The TESCREAL bundle: Eugenics and the promise of utopia through artificial general intelligence | First Monday](#)
- [ACL Is Not an AI Conference](#)
- [ACL is not an AI Conference \(?\)](#)
- [Natural Language Processing RELIES on Linguistics](#)
- [The Ultimate Guide to Word Embeddings](#)
- [Neural Networks, Manifolds, and Topology -- colah's blog](#)
- [RLHF: Reinforcement Learning from Human Feedback](#)
- [Language Processing Pipelines · spaCy Usage Documentation](#)

8/21: Machine learning foundations: Logistic regression

- [Logistic Regression](#) (Section 5.2) in Speech and Language Processing (3rd ed. draft)

8/26: Tokenizer; Morphology

- [Word and Subword Tokenization \(Section 2.5\)](#) in Speech and Language Processing (3rd ed. draft)
- [Neural Machine Translation of Rare Words with Subword Units](#) (the first paper that applied BPE for tokenization)
- ['Let's build the GPT Tokenizer' by Andrej Karpathy](#) (practical tour of BPE with a focus on LLMs)
- [Between words and characters: A Brief History of Open-Vocabulary Modeling and Tokenization in NLP](#) (an excellent survey on tokenization)

8/28: Neural networks foundations: Feedforward neural networks

- [Logistic Regression](#) (Section 5.3) in Speech and Language Processing (3rd ed. draft)
- [Neural Networks](#) (until Section 7.4) in Speech and Language Processing (3rd ed. draft)
- [Neural Networks, Manifolds, and Topology -- colah's blog](#)

9/4: Vector Semantics and Embeddings

- [Vector Semantics and Embeddings](#) (Sections 6.1, 6.2., 6.4, 6.8+) in Speech and Language Processing (3rd ed. draft)
- [\[1310.4546\] Distributed Representations of Words and Phrases and their Compositionality](#) (the paper that introduced skip-gram)
- <https://fasttext.cc/docs/en/unsupervised-tutorial.html> (check out this tutorial for how to train your own word embeddings easily)

- [Understanding and Creating Word Embeddings | Programming Historian](#) (see how historians can use word embeddings)

9/9: Vector Semantics and Embeddings (Continued)

- [Jurafsky and Martin: Speech and Language Processing](#) Section 7.4

9/11: Language Modeling

- [Jurafsky and Martin: Speech and Language Processing](#) Sections 3.1-3.5; Section 7.6 (Feedforward Neural Language Modeling); Chapter 8 (RNNs & LSTMs)

9/16: Machine Translation: Seq2Seq

- [Jurafsky and Martin: Speech and Language Processing](#) Section 13.2

9/18: Machine Translation: BLEU, Decoding, Attention

- Attention ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Section 9.1
- Decoding ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Section 10.2
- MT Evaluation ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Section 13.6

9/23: Transformer

- [Jurafsky and Martin: Speech and Language Processing](#) Chapter 9
- [Attention Is All You Need](#) - The paper where the transformer architecture is introduced
- [The Illustrated Transformer – Jay Alammar](#) - Helpful illustration
- [\[Public. Approved\] Intro to Transformers](#) - Slides by Lucas Beyer

9/25: Pretraining & Finetuning

- Pretraining with MLM & finetuning encoder-only MLM-pretrained models ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Chapter 11
- LM pretraining & finetuning ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Section 10.3

9/25 (cont.): Pretraining Data

- [Jurafsky and Martin: Speech and Language Processing](#) Section 10.3.2

10/14: Prompting

- [Jurafsky and Martin: Speech and Language Processing](#) Chapter 12
- [The Prompt Report: A Systematic Survey of Prompting Techniques](#) - Useful survey/overview

10/16: RLHF

- The original work where RLHF is proposed: [Deep reinforcement learning from human preferences](#)
- Using RLHF to train a summarization model using human feedback: [Learning to summarize from human feedback](#)
- This work is commonly referred to as InstructGPT; the first work to bring RLHF for further adjusting large language models and what lead to ChatGPT: [Training language models to follow instructions with human feedback](#)

Additional blogs that can help to understand:

- [Illustrating Reinforcement Learning from Human Feedback \(RLHF\)](#)
- [RLHF: Reinforcement Learning from Human Feedback](#)

10/21: Transformer types & Practical considerations

- Decoder-only transformer ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Figure 9.15 in Section 9.5
- Classification head / classifier head ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Figure 11.9 in Section 11.4.1 (illustrated with encoder-only transformer)
- Encoder-only transformer ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Section 11.1
 - “Remember that we said the models of Chapter 9 are sometimes called decoderonly, because they correspond to the decoder part of the encoder-decoder model we will introduce in Chapter 13. By contrast, the masked language models of this chapter are sometimes called encoder-only, because they produce an encoding for each input token but generally aren’t used to produce running text by decoding/sampling. That’s an important point: masked language models are not used for generation. They are generally instead used for interpretative tasks.”
- Encoder-decoder transformer ⇒ [Jurafsky and Martin: Speech and Language Processing](#) Figure 13.5 and 13.6 in Section 13.3
- Hyung Won Chung’s lecture ([recording](#), [slides](#))

10/23: Weight and Key-value Cache Quantization

- KV cache original paper: [Efficiently Scaling Transformer Inference](#)
- State-of-the-art KV cache quantization: [KIVI: A Tuning-Free Asymmetric 2bit Quantization for KV Cache](#)

Additional blogs that can help to understand:

- [A Gentle Introduction to 8-bit Matrix Multiplication for transformers at scale using Hugging Face Transformers, Accelerate and bitsandbytes](#)
- [Unlocking Longer Generation with Key-Value Cache Quantization](#)
- [Understanding Quantization for LLMs | by LM Po | Medium](#)

- [Transformers KV Caching Explained | by João Lages | Medium](#)

10/28: (Q)LoRA

- [Jurafsky and Martin: Speech and Language Processing](#) Section 10.5.3 (PEFT & LoRA)
- LORA original paper: [LoRA: Low-Rank Adaptation of Large Language Models](#)
- QLoRA original paper: [QLoRA: Efficient Finetuning of Quantized LLMs](#)
- PEFT survey: [Parameter-Efficient Fine-Tuning for Large Models: A Comprehensive Survey](#)

Additional blogs that can help to understand:

- [Making LLMs even more accessible with bitsandbytes, 4-bit quantization and QLoRA](#)
- [QLoRA: Fine-Tuning Large Language Models \(LLM's\)](#)

10/28 (cont.): QA landscape

- [QA Dataset Explosion: A Taxonomy of NLP Resources for Question Answering and Reading Comprehension](#)

10/30: Dense retrieval; Answer extraction

- [Jurafsky and Martin: Speech and Language Processing](#) Sections 14.1–14.2
- [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#) Section 4.2 for answer extraction

Additional blogs that can help to understand:

- [How to Build an Open-Domain Question Answering System? | Lil'Log](#)

11/4: Retrieval augmented generation (RAG); Summarization

- [ACL 2023 Tutorial: Retrieval-based Language Models and Applications](#) a comprehensive tutorial
-